

Sparse Attention Routing for Low-Latency Retrieval

Anonymous submission · under review

We present a routing scheme that reduces attention cost on long retrieval contexts by selecting a sparse subset of key positions before the softmax. On three public benchmarks the method matches dense attention within 0.4 points while cutting wall-clock latency by 38 percent at a 32k context length.